

## 1. AI for the game of chess

### MDP:

- Set of **States S**:  $S = \{0, \dots, 12\}^{64}$ , where each entry  $s_i$  of  $s \in S$  represents a tile ( $8 \cdot 8 = 64$  tiles) and its value defines what kind of piece occupies it. Values of 0 indicate an empty field, 1 to 6 the white chess pieces and 7 to 12 the black chess pieces. For example, if  $s_5 = 1$ , there is a white pawn on tile A5 and if  $s_{13} = 8$ , there is a black knight on tile B5 ( $13 = 5 + 8$ ).
- Set of **Actions A**:  $A = \{(i, j) \mid i, j \in \{1, \dots, 64\}, i \neq j\}$  - an action is a pair of board coordinates  $i$  and  $j$ , that describes the action of moving any piece from position  $i$  to position  $j$ .
- **Probabilistic State Dynamics**  $p(s'|s, a, \text{policy of the opponent})$ :
  - Any **illegal move** (e.g. selecting a tile without a chess piece on it, selecting a tile with a chess piece of the opposing color, moving a pawn sideways etc.) is mapped back onto the same state  $s$  with a probability of 1 and the agent needs to pick a action again.
  - Any **legal move** results deterministically in a board configuration  $\hat{s} \in S$ . Depending on the opponent's policy  $\pi^o(\cdot | \hat{s})$ , we obtain a new state  $s' \in S$  with probability  $\pi^o(s' | \hat{s})$ .
- **Reward Dynamics**  $p(R_{t+1} | s, a)$ :
  - Any **illegal move** will be penalized by e.g. -10
  - Any action resulting in **losing a chess piece** due to the opponent's action will give a negative reward depending on the role of the taken piece (queen gives higher penalty than a pawn) e.g. between -5 and -20
  - Any action that **takes an opponents chess piece** will get a reward dependent on the role of the taken piece (queen gives higher reward than a pawn) e.g. between +5 and +20
  - **Taking the opponent's king** or setting the opponent **checkmate** will win the game and yields a high reward, e.g. +100.
  - Checking the opponent will yield a positive reward e.g +10
  - By defining the rewards as above the agent is encouraged to learn what moves are actually allowed by himself, additionally to learning a good strategy.
- **Policy**  $\pi(a|s)$ :  $\pi(a|s) = (p_{a1}, p_{a2}, \dots, p_{a4032})$ , where  $\sum_{\text{over all } p} = 1, \forall s \in S, a \in A$

## 2. LunarLander MDP & Policy

MDP:

- Set of **States S**:  $S = \{(x, y, rot, v_x, v_y, v_{rot}, leftLegContact, rightLegContact) \mid x \in \mathbb{Z}, y \in \mathbb{Z}, rot \in [0, 2 * \pi), v_x, v_y, v_{rot} \in \mathbb{R}, leftLegContact, rightLegContact \in \{0, 1\}\}$
- Set of **Actions A**: (do nothing, fire left orientation engine, fire main engine, fire right orientation engine)
- **Probabilistic State Dynamics**  $p(s'|s,a)$  = Since we have infinite fuel, no action will fail. Therefore, all actions will lead to a follow up state  $s'$  directly following the physics of the environment with  $p(s'|s,a)=1$  and  $p(t|s,a)=0$  for all other states  $t$ .
- **Reward Dynamics**  $r(s,a)$  = [moving down ~ +100-140, zero speed ~ +100-140, moving away from landing pad = losing reward, coming to rest = +100, crashing = -100, firing engine = -0.3 per frame, solved = +200, leg ground contact = +10 per leg]
- **Policy**  $\pi(a|s)$ :  
 $\pi(a|s) = (p_{do-nothing}, p_{fire-left-engine}, p_{fire-right-engine}, p_{main-engine})$ ,  
where  $\sum_{over\ all\ p} = 1, s \in S, a \in A$

### 3. Discuss the Policy Evaluation and Policy Iteration algorithms

Explain what the environment dynamics (i.e. reward function and state transition function) are and give at least two examples.

- State transition function: Provides the probability of arriving in state  $s'$  when performing action  $a$  in state  $s$   $p(s'|s, a)$ .
  - E.g. in a self-coded grid world the probability of getting to the next state can be coded as 1, so there is no uncertainty about  $s'$ , i.e. when the agent wants to move up from (0,0) it will always get to position (1,0).
  - E.g. in a 2-player game like go or chess, where the opponent's moves cannot be predicted, the future game state is probabilistic, therefore the state-transition-function is probabilistic and usually unknown.
- Reward function  $r(s,a)$ : Provides the expected reward given the state  $s$  and action  $a$ .
  - Again, in the self-coded grid world, the developer defines which states give which rewards, e.g. state (4,4) is a goal state and transitioning into it yields a reward of +100 and all other transitions might get penalized by a small amount of -1 to encourage taking the shortest route as possible.
  - In a 2-player game like go or chess, where the opponent's moves cannot be predicted, an action may result in a better or worse outcome depending on the opponent's reaction. Therefore, reward is probabilistic and the reward function returns the expected (mean) reward.

Discuss: Are the environment dynamics generally known and can practically be used to solve a problem with RL?

- If we have a controlled environment that is e.g. self-coded (like Gridworld), then probably yes.
- In an uncontrolled environment (e.g. letting a robot interact with the real world), having these is unlikely and can only be observed through interaction with the environment, i.e. an estimate can be obtained.